

## ▼ Trabajo Final - Saul Esquivel Condori

Importamos los módulos necesarios

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import matplotlib as mpl
from datetime import datetime
plt.rcParams.update({'figure.figsize': (10, 7), 'figure.dpi': 120})
from statsmodels.tsa.seasonal import seasonal_decompose
from statsmodels.tsa.api import ExponentialSmoothing, SimpleExpSmoothing, Holt

from google.colab import files

from google.colab import drive
drive.mount('/content/drive')

↳ Drive already mounted at /content/drive; to attempt to forcibly remount, call drive.n
```

Leemos nuestra serie de tiempo desde un archivo csv

```
# Importar datos en formato csv.
#custom_date_parser = lambda x: datetime.strptime(x, "%d/%m/%Y")
Salesd = pd.read_csv('/content/drive/MyDrive/Forecast/SalesDemand.csv',encoding='latin-1',

Salesd['date'] = pd.to_datetime(Salesd['date'], format='%d/%m/%Y')
#Salesd['periodo'] = Salesd['date'].dt.strftime('%Y%m')

#Validamos la correcta carga de la base de datos
=====
Salesd.head(10)
```

```

      date  store  item  sales
0  2013-01-01    1.0    1.0   13.0
1  2013-01-02    1.0    1.0   11.0
2  2013-01-03    1.0    1.0   14.0
3  2013-01-04    1.0    1.0   13.0

```

```
#list(Salesd.groupby('item'))
```

```

5  2013-01-06    1.0    1.0   12.0

```

```
Salesd.dtypes
```

	date	store	item	sales
0	datetime64[ns]	float64	float64	float64
1				
2				
3				
4				
5				
6				
7				
8				
9				
10				
11				
12				
13				
14				
15				
16				
17				
18				
19				
20				
21				
22				
23				
24				
25				
26				
27				
28				
29				
30				
31				
32				
33				
34				
35				
36				
37				
38				
39				
40				
41				
42				
43				
44				
45				
46				
47				
48				
49				
50				
51				
52				
53				
54				
55				
56				
57				
58				
59				
60				
61				
62				
63				
64				
65				
66				
67				
68				
69				
70				
71				
72				
73				
74				
75				
76				
77				
78				
79				
80				
81				
82				
83				
84				
85				
86				
87				
88				
89				
90				
91				
92				
93				
94				
95				
96				
97				
98				
99				
100				
101				
102				
103				
104				
105				
106				
107				
108				
109				
110				
111				
112				
113				
114				
115				
116				
117				
118				
119				
120				
121				
122				
123				
124				
125				
126				
127				
128				
129				
130				
131				
132				
133				
134				
135				
136				
137				
138				
139				
140				
141				
142				
143				
144				
145				
146				
147				
148				
149				
150				
151				
152				
153				
154				
155				
156				
157				
158				
159				
160				
161				
162				
163				
164				
165				
166				
167				
168				
169				
170				
171				
172				
173				
174				
175				
176				
177				
178				
179				
180				
181				
182				
183				
184				
185				
186				
187				
188				
189				
190				
191				
192				
193				
194				
195				
196				
197				
198				
199				
200				
201				
202				
203				
204				
205				
206				
207				
208				
209				
210				
211				
212				
213				
214				
215				
216				
217				
218				
219				
220				
221				
222				
223				
224				
225				
226				
227				
228				
229				
230				
231				
232				
233				
234				
235				
236				
237				
238				
239				
240				
241				
242				
243				
244				
245				
246				
247				
248				
249				
250				
251				
252				
253				
254				
255				
256				
257				
258				
259				
260				
261				
262				
263				
264				
265				
266				
267				
268				
269				
270				
271				
272				
273				
274				
275				
276				
277				
278				
279				
280				
281				
282				
283				
284				
285				
286				
287				
288				
289				
290				
291				
292				
293				
294				
295				
296				
297				
298				
299				
300				
301				
302				
303				
304				
305				
306				
307				
308				
309				
310				
311				
312				
313				
314				
315				
316				
317				
318				
319				
320				
321				
322				
323				
324				
325				
326				
327				
328				
329				
330				
331				
332				
333				
334				
335				
336				
337				
338				
339				
340				
341				
342				
343				
344				
345				
346				
347				
348				
349				
350				
351				
352				
353				
354				
355				
356				
357				
358				
359				
360				
361				
362				
363				
364				
365				
366				
367				
368				
369				
370				
371				
372				

```
--  -----
 0  date    913000 non-null  datetime64[ns]
 1  store   913000 non-null  float64
 2  item    913000 non-null  float64
 3  sales   913000 non-null  float64
dtypes: datetime64[ns](1), float64(3)
memory usage: 34.8 MB
```

Sales.head()

	date	store	item	sales	🔗
0	2013-01-01	1.0	1.0	13.0	
1	2013-01-02	1.0	1.0	11.0	
2	2013-01-03	1.0	1.0	14.0	
3	2013-01-04	1.0	1.0	13.0	
4	2013-01-05	1.0	1.0	10.0	

```
Sales_store = Sales[['store','sales']]
Sales_store_2 = Sales_store.groupby(['store']).sum()
Sales_store_2.sort_values(by="sales", ascending=False)
#tomaremos las 5 tiendas con mayores ventas totales y que representan mas del 50% de venta
```

	sales	🔗
store		
<b>2.0</b>	6120128.0	
<b>8.0</b>	5856169.0	
<b>3.0</b>	5435144.0	
<b>10.0</b>	5360158.0	
<b>9.0</b>	5025976.0	
<b>4.0</b>	5012639.0	
<b>1.0</b>	4315603.0	
<b>5.0</b>	3631016.0	
<b>6.0</b>	3627670.0	
<b>7.0</b>	3320009.0	

```
Sales_item = Sales[['item','sales']]
Sales_item_2 = Sales_item.groupby(['item']).sum()
Sales_item_2.sort_values(by="sales", ascending=False)
#tomaremos los 18 items con mayores ventas totales y que representan mas del 50% de ventas
```

sales 

item	sales
<b>15.0</b>	1607442.0
<b>28.0</b>	1604713.0
<b>13.0</b>	1539621.0
<b>18.0</b>	1538876.0
<b>25.0</b>	1473334.0
<b>45.0</b>	1471467.0
<b>38.0</b>	1470330.0
<b>22.0</b>	1469971.0
<b>36.0</b>	1406548.0
<b>8.0</b>	1405108.0
<b>10.0</b>	1337133.0
<b>11.0</b>	1271925.0
<b>12.0</b>	1271534.0
<b>29.0</b>	1271240.0
<b>33.0</b>	1270183.0
<b>24.0</b>	1205975.0
<b>50.0</b>	1203009.0
<b>35.0</b>	1201541.0
<b>14.0</b>	1071531.0
<b>31.0</b>	1070845.0
<b>46.0</b>	1070764.0
<b>2.0</b>	1069564.0
<b>7.0</b>	1068777.0
<b>6.0</b>	1068281.0
<b>9.0</b>	938379.0
<b>48.0</b>	937703.0
<b>43.0</b>	936635.0
<b>26.0</b>	869981.0
<b>20.0</b>	867641.0
<b>32.0</b>	803107.0
<b>39.0</b>	801311.0

**19.0** 736892 0

```
Sales1 = Sales[Sales['store'].isin([2,8,3,10,9])]
Sales2 = Sales1[Sales1['item'].isin([15,28,13,18,25,45,38,22,36,8,10,11,12,29,33,24,50,35])]
```

**21.0** /36190.0

```
Sales2.info()
# el analisis lo haremos sobre 73'040 registros
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 164340 entries, 129646 to 912999
Data columns (total 4 columns):
 #   Column  Non-Null Count  Dtype  
---  -- 
 0   date    164340 non-null   datetime64[ns]
 1   store   164340 non-null   float64 
 2   item    164340 non-null   float64 
 3   sales   164340 non-null   float64 
dtypes: datetime64[ns](1), float64(3)
memory usage: 6.3 MB
```

**34.0** 469935.0

```
#Sales2[['date','sales']]
Sales3 = Sales2[Sales2['date']=='2013-01-01']
Sales3.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 90 entries, 129646 to 911174
Data columns (total 4 columns):
 #   Column  Non-Null Count  Dtype  
---  -- 
 0   date    90 non-null    datetime64[ns]
 1   store   90 non-null    float64 
 2   item    90 non-null    float64 
 3   sales   90 non-null    float64 
dtypes: datetime64[ns](1), float64(3)
memory usage: 3.5 KB
```

```
Sales3 = Sales2[['date','sales']]
```

```
Sales4 = Sales3.groupby(['date']).sum()
Sales4.head()
# se observa que falta el index, que resetearemos a continuación.
```

	sales
<b>date</b>	
<b>2013-01-01</b>	4133.0
<b>2013-01-02</b>	4255.0
<b>2013-01-03</b>	4308.0
<b>2013-01-04</b>	4636.0
<b>2013-01-05</b>	5033.0

```
Sales4.reset_index(level=0, inplace=True)
#Sales4['Id'] = np.arange(len(Sales4)) sirve para agregar columna con numeros consecutivos
Sales4
```

	date	sales	edit
0	2013-01-01	4133.0	
1	2013-01-02	4255.0	
2	2013-01-03	4308.0	
3	2013-01-04	4636.0	
4	2013-01-05	5033.0	
...	...	...	
1821	2017-12-27	6137.0	
1822	2017-12-28	6595.0	
1823	2017-12-29	7223.0	
1824	2017-12-30	7714.0	
1825	2017-12-31	8091.0	

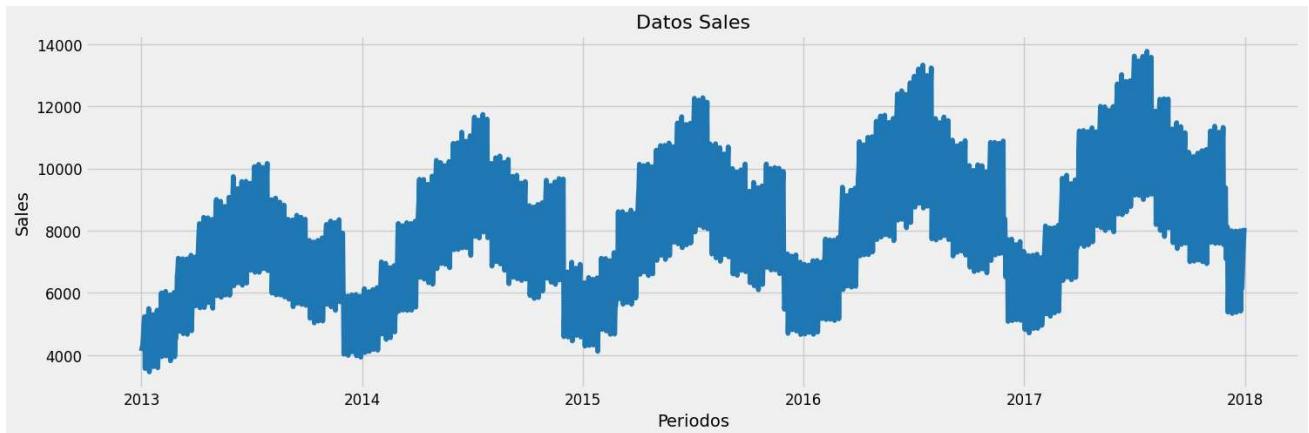
1826 rows × 2 columns

```
Sales4.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1826 entries, 0 to 1825
Data columns (total 2 columns):
 #   Column   Non-Null Count   Dtype  
--- 
 0   date      1826 non-null    datetime64[ns]
 1   sales     1826 non-null    float64 
dtypes: datetime64[ns](1), float64(1)
memory usage: 28.7 KB
```

## Graficamos la serie de tiempo

```
plt.figure(figsize=(16,5), dpi=100)
plt.plot(Sales4.date, Sales4.sales,color='tab:blue')
plt.xlabel('Periodos')
plt.ylabel('Sales')
plt.title('Datos Sales', fontsize=16)
#plt.gca().set(title='Serie CQI', xlabel=df_1.RESULT_TIME, ylabel=df_1.CQI)
plt.show()
```



```

Sales4['Year'] = Sales4['date'].dt.strftime('%Y')
Sales4['Month'] = Sales4['date'].dt.strftime('%m')
Sales4['Month'] = Sales4.Month.astype(int)
Sales4.head()
Sales4.info()

```

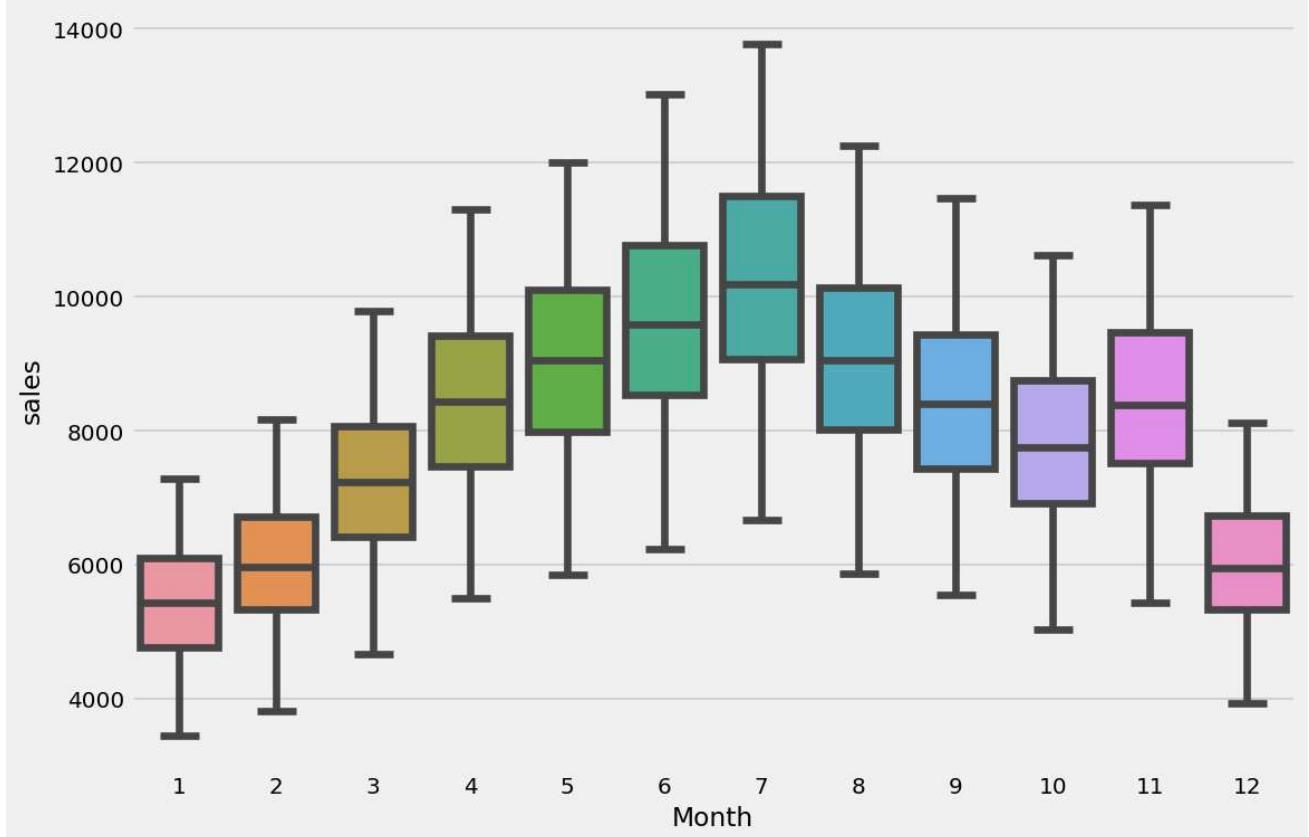
```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1826 entries, 0 to 1825
Data columns (total 4 columns):
 #   Column   Non-Null Count   Dtype  
--- 
 0   date      1826 non-null    datetime64[ns]
 1   sales     1826 non-null    float64 
 2   Year       1826 non-null    object  
 3   Month      1826 non-null    int64  
dtypes: datetime64[ns](1), float64(1), int64(1), object(1)
memory usage: 57.2+ KB

```

```
sns.boxplot(x = 'Month', y = 'sales', data = Sales4)
```

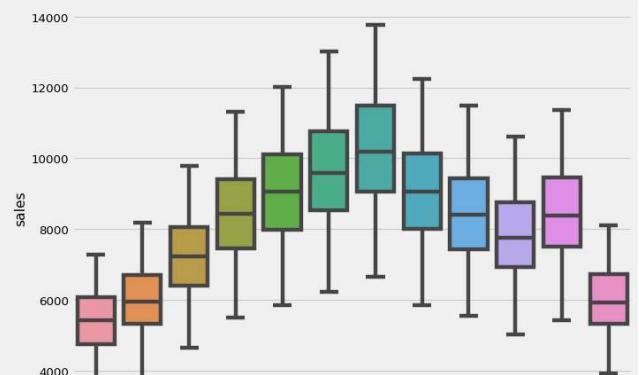
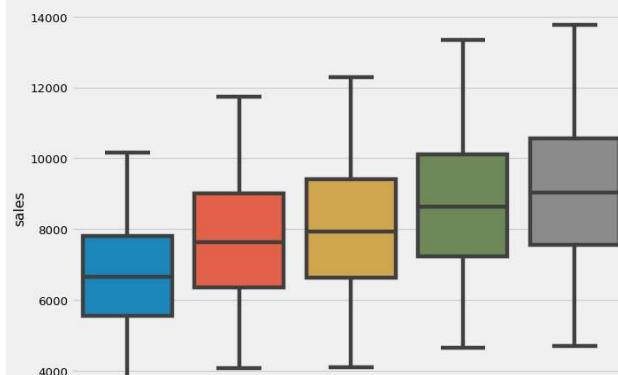
```
<Axes: xlabel='Month', ylabel='sales'>
```



## Graficamos las cajas agregadas

```
# Draw Plot
fig, axes = plt.subplots(1, 2, figsize=(20,7), dpi= 80)
sns.boxplot(x='Year', y='sales', data=Sales4, ax=axes[0])
sns.boxplot(x='Month', y='sales', data=Sales4, ax=axes[1])
```

<Axes: xlabel='Month', ylabel='sales'>



## Agrupando periodos

```
# agrupando por periodos específicos
Sales_mes= Sales4[['date','sales']]
```

```
Sales_mes.set_index('date',inplace = True)
```

```
Sales_mes.head(10)
```

	sales	date
2013-01-01	4133.0	
2013-01-02	4255.0	
2013-01-03	4308.0	
2013-01-04	4636.0	
2013-01-05	5033.0	
2013-01-06	5235.0	
2013-01-07	3567.0	
2013-01-08	4224.0	
2013-01-09	4217.0	
2013-01-10	4375.0	

```
# agrupamos los datos
Sales_mes = Sales_mes.resample('M').sum()
Sales_year = Sales_mes.resample('Y').mean()
Sales_quarter = Sales_mes.resample('Q').mean()
```

```
Sales_mes.head()
```

sales 

date

2013-01-31	138832.0
2013-02-28	140085.0
2013-03-31	189060.0
2013-04-30	208446.0

```
print(Sales_mes.shape)
print(Sales_year.shape)
print(Sales_quarter.shape)
```

```
(60, 1)
(5, 1)
(20, 1)
```

```
Sales_mes.head(10)
```

sales 

date

2013-01-31	138832.0
2013-02-28	140085.0
2013-03-31	189060.0
2013-04-30	208446.0
2013-05-31	233266.0
2013-06-30	243439.0
2013-07-31	261151.0
2013-08-31	234651.0
2013-09-30	210757.0
2013-10-31	200192.0

```
Sales_quarter.head(10)
```

**sales** 

**date**

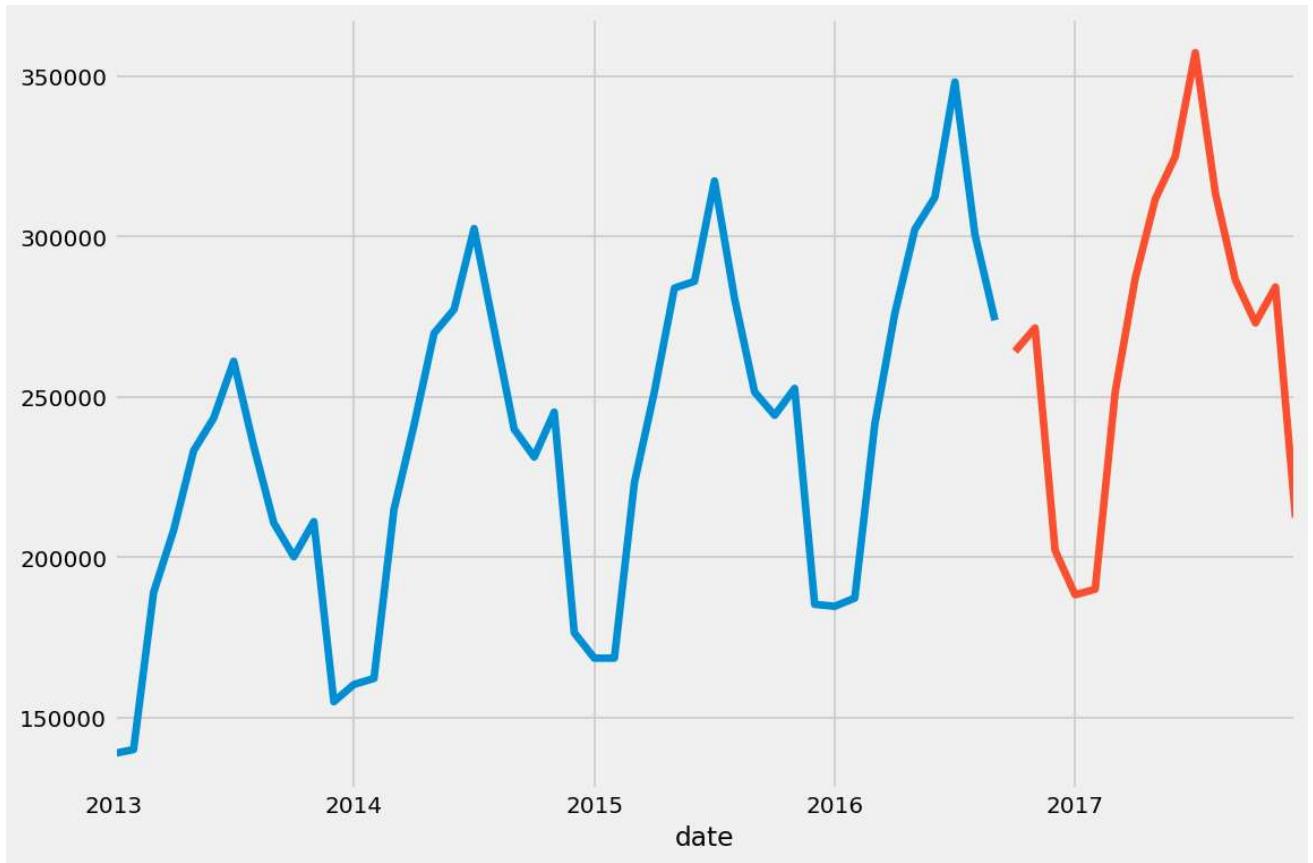
date	sales
2013-03-31	155992.333333
2013-06-30	228383.666667
2013-09-30	235519.666667
2013-12-31	188785.333333
2014-03-31	179207.333333
2014-06-30	262778.666667
2014-09-30	271198.666667
2014-12-31	217692.000000

```
plt.figure(figsize=(16,5), dpi=100)
plt.plot(Sales_mes.sales,color='tab:green',label='Mensual')
plt.plot(Sales_year.sales,color='tab:blue',label='Anual')
plt.plot(Sales_quarter.sales,color='tab:red',label='Trimestral')
plt.legend()
plt.show()
```

**Prophet**

```
#divide into train and validation set
train = Sales_mes[:int(0.75*(len(Sales_mes)))]
valid = Sales_mes[int(0.75*(len(Sales_mes))):]
```

```
#plotting the data  
train['sales'].plot()  
valid['sales'].plot()  
plt.show()
```



```
train.head()
```

**sales** 

**date**


---

**2013-01-31** 138832.0

**2013-02-28** 140085.0

```
print(train.shape)
```

```
print(valid.shape)
```

```
(45, 1)  
(15, 1)
```

```
train_prophet = pd.DataFrame()  
train_prophet['ds'] = train.index  
train_prophet['y'] = train.sales.values
```

```
train_prophet.head()
```

	<b>ds</b>	<b>y</b>	
<b>0</b>	2013-01-31	138832.0	
<b>1</b>	2013-02-28	140085.0	
<b>2</b>	2013-03-31	189060.0	
<b>3</b>	2013-04-30	208446.0	
<b>4</b>	2013-05-31	233266.0	

```
import os  
# Let cmdstanpy know where CmdStan is  
os.environ["CMDSTAN"] = "./cmdstan-2.23.0"
```

```
from prophet import Prophet
```

```
#instantiate Prophet with only yearly seasonality as our data is monthly  
model = Prophet()  
model.fit(train_prophet) #fit the model with your dataframe
```

```
INFO:prophet:Disabling weekly seasonality. Run prophet with weekly_seasonality=True to  
INFO:prophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to  
DEBUG:cmdstanpy:input tempfile: /tmp/tmpivo66ed6/xhraar48.json  
DEBUG:cmdstanpy:input tempfile: /tmp/tmpivo66ed6/_160n6tg.json  
DEBUG:cmdstanpy:idx 0  
DEBUG:cmdstanpy:running CmdStan, num_threads: None  
DEBUG:cmdstanpy:CmdStan args: ['/usr/local/lib/python3.9/dist-packages/prophet/stan_n  
04:06:19 - cmdstanpy - INFO - Chain [1] start processing  
INFO:cmdstanpy:Chain [1] start processing  
04:06:20 - cmdstanpy - INFO - Chain [1] done processing  
INFO:cmdstanpy:Chain [1] done processing  
<prophet.forecaster.Prophet at 0x7fabd1e91670>
```

```
# predict for five months in the future and MS - month start is the frequency
future = model.make_future_dataframe(periods = 15, freq = 'MS')
future.head()
```

ds



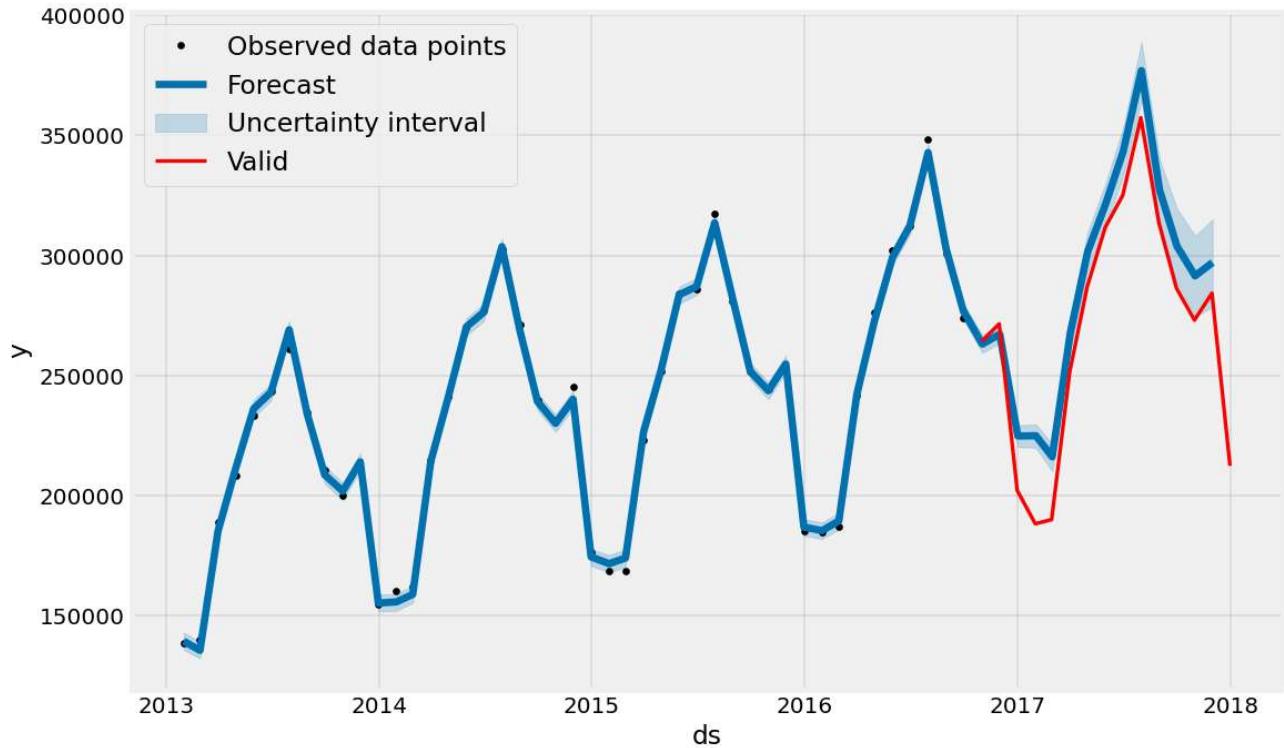
	ds
0	2013-01-31
1	2013-02-28
2	2013-03-31
3	2013-04-30
4	2013-05-31

```
# now lets make the forecasts
forecast = model.predict(future)
forecast[['ds', 'yhat']]
```

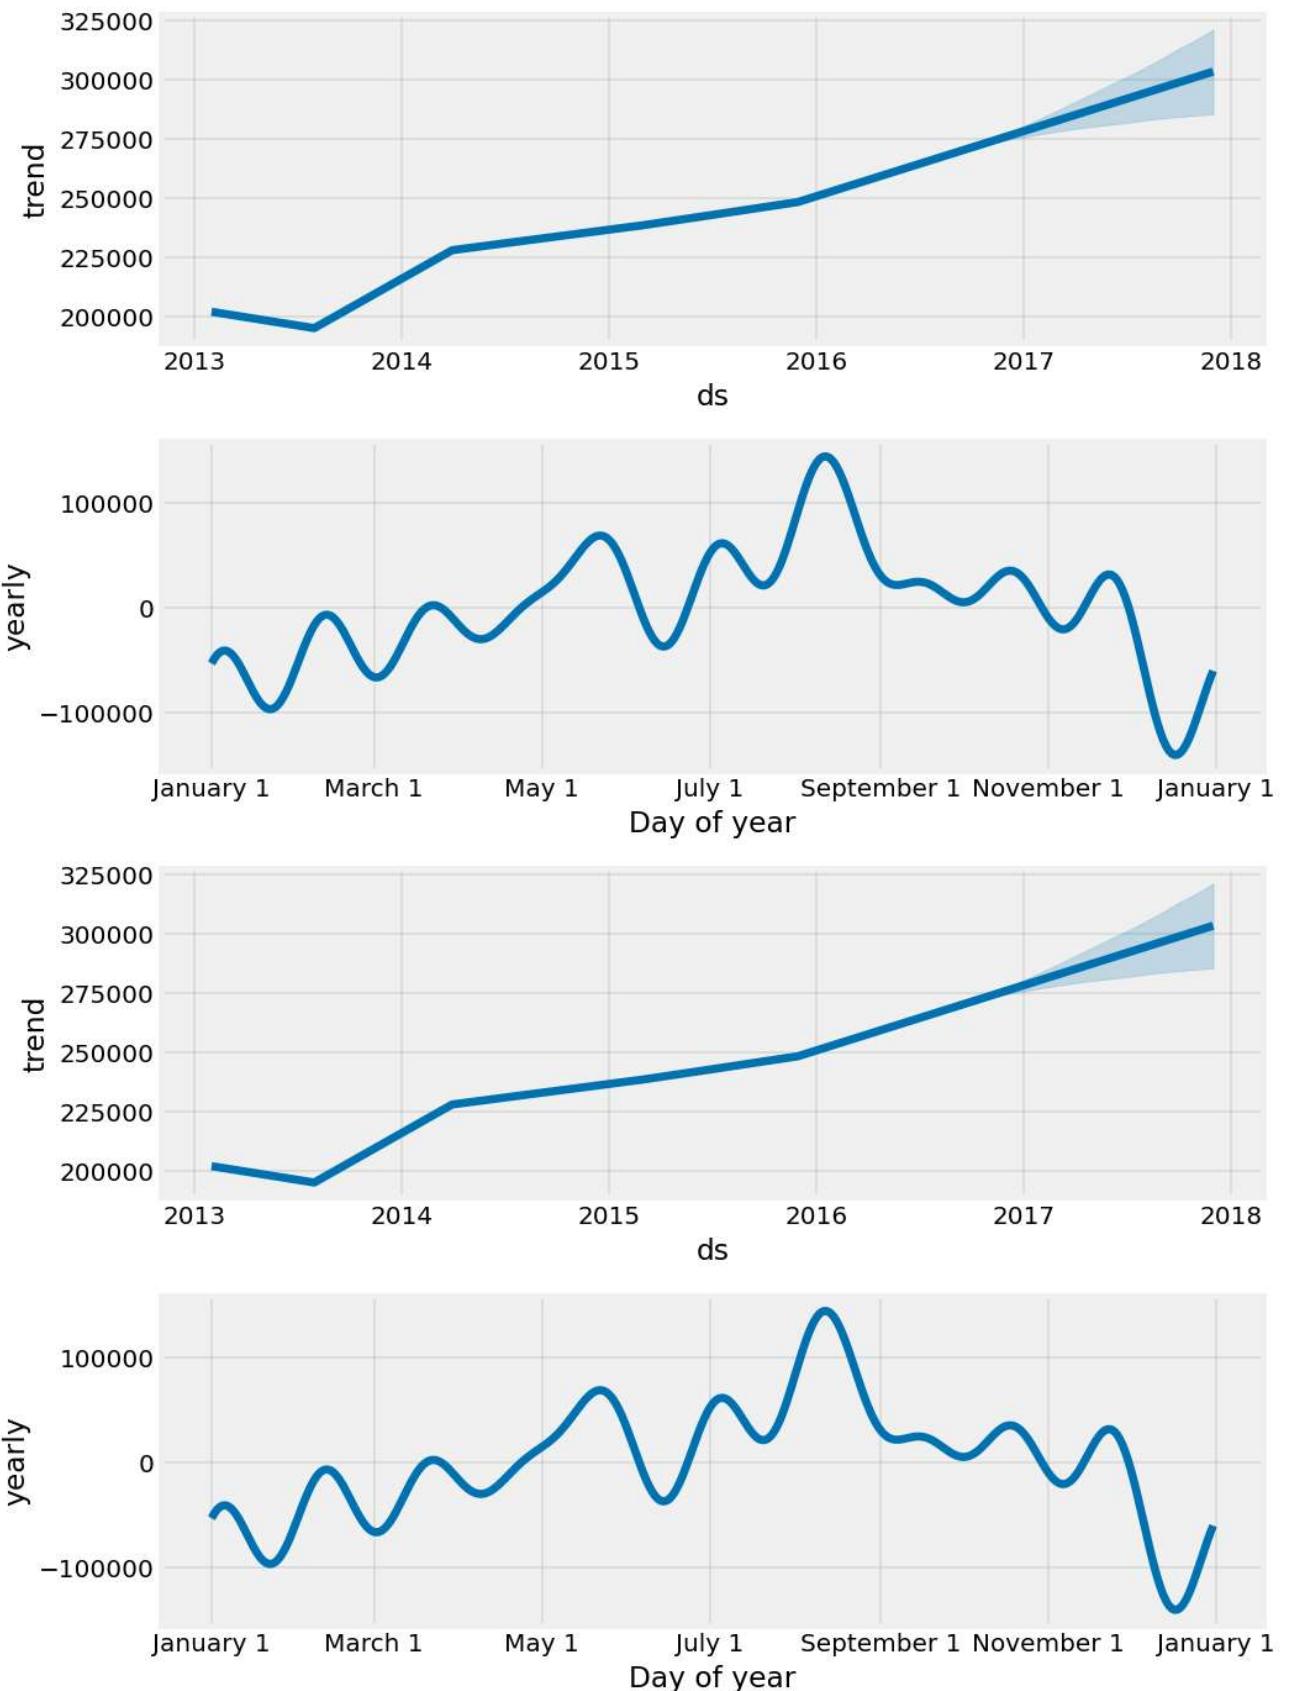
	ds	yhat	edit
0	2013-01-31	139497.884045	
1	2013-02-28	135657.692817	
2	2013-03-31	185367.435162	
3	2013-04-30	211182.155410	
4	2013-05-31	236182.728521	
5	2013-06-30	243149.344005	
6	2013-07-31	269092.539389	
7	2013-08-31	233979.690346	
8	2013-09-30	208702.383870	
9	2013-10-31	201909.467924	
10	2013-11-30	214109.297663	
11	2013-12-31	155329.494634	
12	2014-01-31	155831.463600	
13	2014-02-28	158974.757649	
14	2014-03-31	214365.819564	
15	2014-04-30	241084.893925	
16	2014-05-31	270288.748775	
17	2014-06-30	276433.475463	
18	2014-07-31	303683.271513	
19	2014-08-31	268563.029415	
20	2014-09-30	239233.558177	
21	2014-10-31	230118.372485	
22	2014-11-30	240186.884288	
23	2014-12-31	174553.507575	
24	2015-01-31	171725.170424	
25	2015-02-28	174106.312360	
26	2015-03-31	226654.928860	
27	2015-04-30	252307.292349	
28	2015-05-31	283663.004499	
29	2015-06-30	287024.215445	
30	2015-07-31	313637.819870	
31	2015-08-31	281915.705756	

```
32 2015_09_30 251781 950573
```

```
fig = model.plot(forecast)
#plot the predictions for validation set
plt.plot(valid, label='Valid', color = 'red', linewidth = 2)
plt.legend()
plt.show()
```



```
model.plot_components(forecast)
```



```
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error, median_abso

def evaluate_forecast(y,pred):
    results = pd.DataFrame({'r2_score':r2_score(y, pred),
                            }, index=[0])
    results['mean_absolute_error'] = mean_absolute_error(y, pred)
    results['median_absolute_error'] = median_absolute_error(y, pred)
    results['mse'] = mean_squared_error(y, pred)
    results['msle'] = mean_squared_log_error(y, pred)
    results['rmse'] = np.sqrt(results['mse'])
    return results

evaluate_forecast(Sales_mes, forecast.yhat)
# se obtiene un r2 = 86%, continuemos buscando otro modelo que otorgue un mejor ajuste.
```

r2_score	mean_absolute_error	median_absolute_error	mse	msle
0.860022	0.370745182	0.226411065	3.8503860408	0.0050141062

```
#definiendo que la estacionalidad sea anual y el modelo multiplicativo
model = Prophet(yearly_seasonality=True,seasonality_mode= 'multiplicative')
model.fit(train_prophet) #fit the model with your dataframe
```

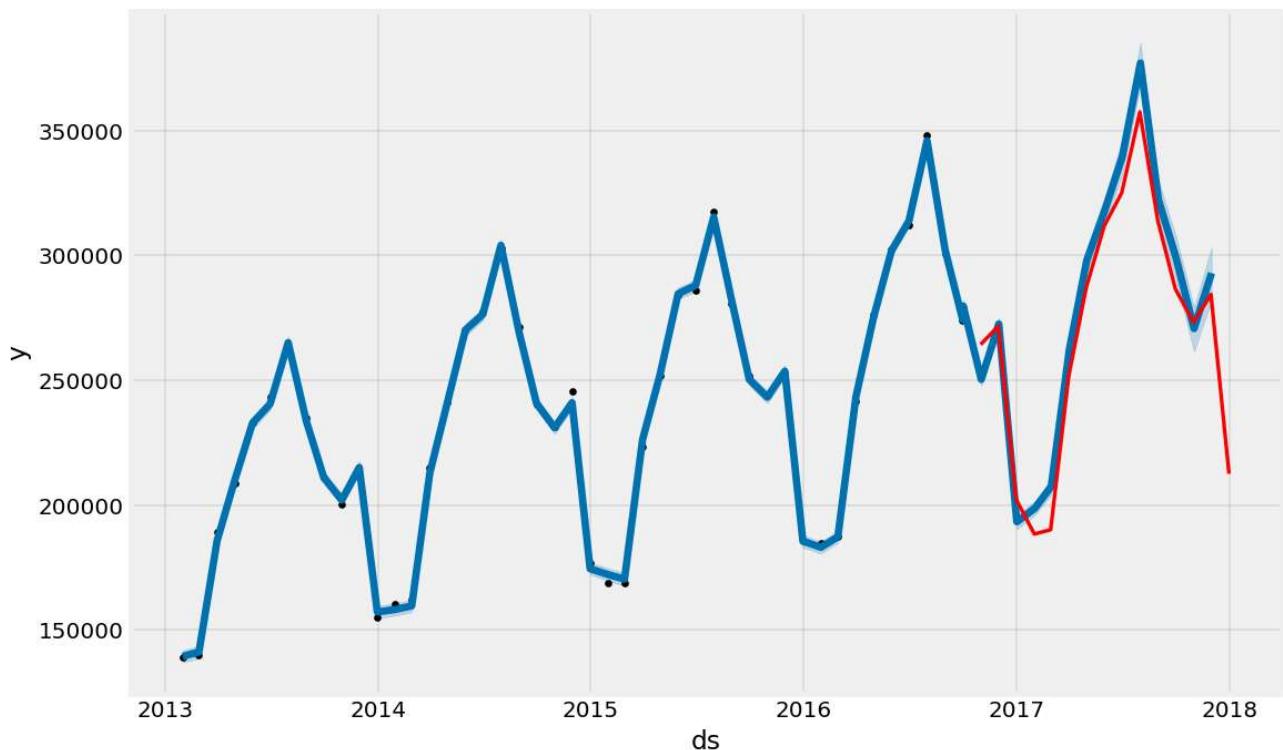
```
INFO:prophet:Disabling weekly seasonality. Run prophet with weekly_seasonality=True to
INFO:prophet:Disabling daily seasonality. Run prophet with daily_seasonality=True to
DEBUG:cmdstanpy:input tempfile: /tmp/tmpivo66ed6/2tmnzzjq.json
DEBUG:cmdstanpy:input tempfile: /tmp/tmpivo66ed6/zczjhah9.json
DEBUG:cmdstanpy:idx 0
DEBUG:cmdstanpy:running CmdStan, num_threads: None
DEBUG:cmdstanpy:CmdStan args: ['/usr/local/lib/python3.9/dist-packages/prophet/stan_n
04:06:22 - cmdstanpy - INFO - Chain [1] start processing
INFO:cmdstanpy:Chain [1] start processing
04:06:23 - cmdstanpy - INFO - Chain [1] done processing
INFO:cmdstanpy:Chain [1] done processing
<prophet.forecaster.Prophet at 0x7fabce1d6160>
```

◀	▶
---	---

```
# predict for five months in the future and MS - month start is the frequency
future = model.make_future_dataframe(periods = 15, freq = 'MS')
forecast = model.predict(future)
fig = model.plot(forecast)
#plot the predictions for validation set

plt.plot(valid, label='Valid', color = 'red', linewidth = 2)

plt.show()
```



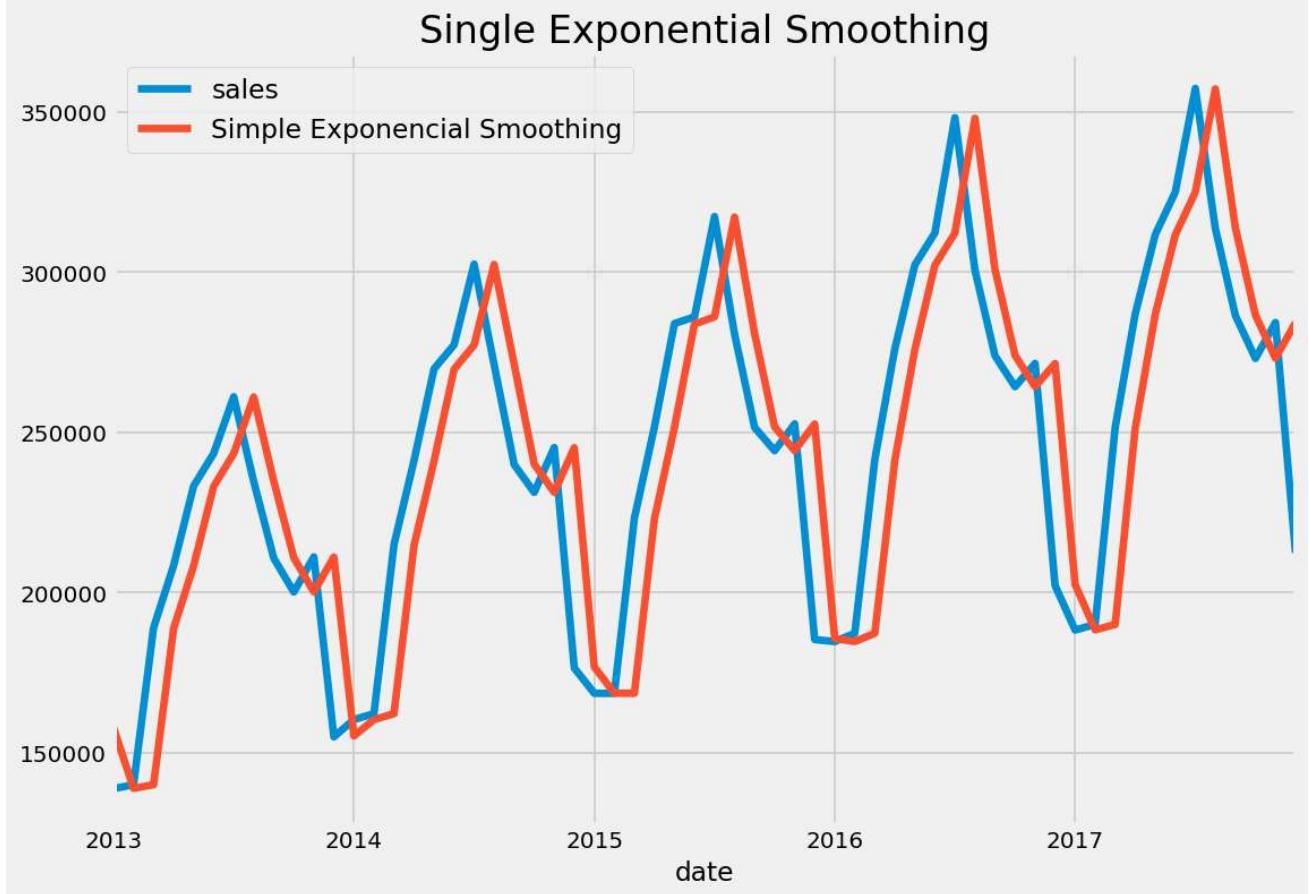
```
evaluate_forecast(Sales_mes, forecast.yhat)
#Se observa que los errores tanto del Prophet Aditivo como Multiplicativo son los mismos,
```

	r2_score	mean_absolute_error	median_absolute_error	mse	msle
0	0.8690003	8626562076	1817768630	3.603366e+08	0.0052841802

## ▼ Simple exponencial smoothing (SES)

```
HWES1 = SimpleExpSmoothing(Sales_mes, initialization_method="estimated").fit()
HWES1_fitted = HWES1.fittedvalues
HWES1_fitted.name = 'Simple Exponential Smoothing'
pd.concat([Sales_mes, HWES1_fitted], axis=1).plot(title='Single Exponential Smoothing')
```

```
<Axes: title={'center': 'Single Exponential Smoothing'}, xlabel='date'>
```



```
results = pd.DataFrame(
    index=["SSE", "AIC", "BIC"])
results["Simple Exponential"] = [HWES1.sse] + [HWES1.aic] + [HWES1.bic]
```

```
print(HWES1.summary())
```

#### SimpleExpSmoothing Model Results

Dep. Variable:	sales	No. Observations:	60
Model:	SimpleExpSmoothing	SSE	62997737631.460
Optimized:	True	AIC	1250.321
Trend:	None	BIC	1254.510
Seasonal:	None	AICC	1251.048
Seasonal Periods:	None	Date:	Sun, 16 Apr 2023
Box-Cox:	False	Time:	04:06:24

Box-Cox Coeff.:	None		
	coeff	code	optimized
smoothing_level	0.9950000	alpha	True
initial_level	1.5788e+05	1.0	True

## ▼ Double exponential smoothing (DES)

```
HWES2_ADD = ExponentialSmoothing(Sales_mes,trend='add').fit()#.fittedvalues
HWES2_MUL = ExponentialSmoothing(Sales_mes,trend='mul').fit()#.fittedvalues
HWES2_ADD_fitted = HWES2_ADD.fittedvalues
HWES2_MUL_fitted = HWES2_MUL.fittedvalues
HWES2_ADD_fitted.name = 'Additive Double Exp. Smoothing'
HWES2_MUL_fitted.name = 'multiplicative Double Exp. Smoothin'
pd.concat([Sales_mes, HWES2_ADD_fitted, HWES2_MUL_fitted], axis=1).plot(title='Double Expo
```

## Double Exponential Smoothing: Additive and Multiplicative Trend



```
results["Double Exp. - Additive"] = [HWES2_ADD.sse] + [HWES2_ADD.aic] + [HWES2_ADD.bic]
results["Double Exp. - Multiplicative"] = [HWES2_MUL.sse] + [HWES2_MUL.aic] + [HWES2_MUL.b
```

```
HWES2_ADD.summary()
```

ExponentialSmoothing Model Results			
Dep. Variable:	sales	No. Observations:	60
Model:	ExponentialSmoothing	SSE	65329777624.011
Optimized:	True	AIC	1256.502
Trend:	Additive	BIC	1264.880
Seasonal:	None	AICC	1258.087
Seasonal Periods:	None	Date:	Sun, 16 Apr 2023
Box-Cox:	False	Time:	04:06:24
Box-Cox Coeff.:	None		
		coeff	code optimized
smoothing_level	0.9950000	alpha	True
smoothing_trend	0.0236905	beta	True
initial_level	1.5788e+05	i.0	True
initial_trend	8746.5879	b.0	True

```
HWES2_MUL.summary()
```

ExponentialSmoothing Model Results			
Dep. Variable:	sales	No. Observations:	60
Model:	ExponentialSmoothing	SSE	69923480990.441
Optimized:	True	AIC	1260.579
Trend:	Multiplicative	BIC	1268.957
Seasonal:	None	AICC	1262.164
Seasonal Periods:	None	Date:	Sun, 16 Apr 2023
Box-Cox:	False	Time:	04:06:24
Box-Cox Coeff.:	None		
		coeff	code optimized
smoothing_level	0.9950000	alpha	True
smoothing_trend	0.0473810	beta	True
initial_level	1.5788e+05	i.0	True
initial_trend	1.0553996	b.0	True

\*Continuamos buscando otro modelo, dado que las tecnicas de suavizamiento vistas no estan considerando estacionalidad

## ▼ Triple exponencial smoothing (TES)

```
HWES3_ADD = ExponentialSmoothing(Sales_mes,trend='add',seasonal='add',seasonal_periods=12)
HWES3_MUL = ExponentialSmoothing(Sales_mes,trend='mul',seasonal='mul',seasonal_periods=12)

HWES3_ADD_fitted = HWES3_ADD.fittedvalues
HWES3_MUL_fitted = HWES3_MUL.fittedvalues
HWES3_ADD_fitted.name = 'Additive TES'
HWES3_MUL_fitted.name = 'multiplicative TES'
pd.concat([Sales_mes, HWES3_ADD_fitted, HWES3_MUL_fitted], axis=1).plot(title='Triple Expo
```

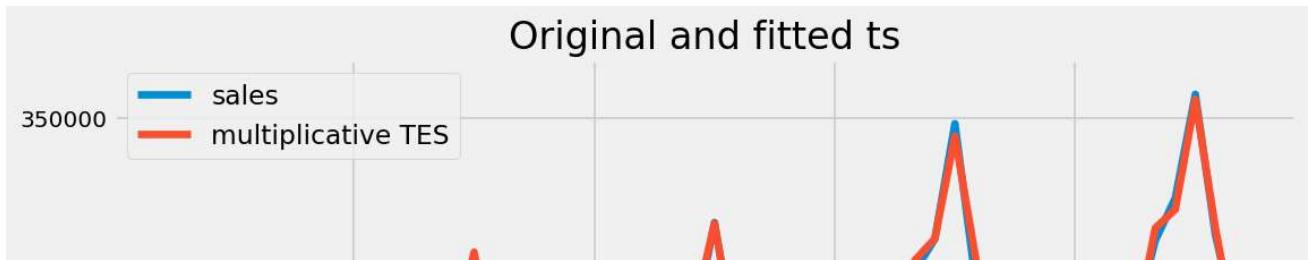
## Triple Exponential Smoothing: Additive and Multiplicative Trend

```
results["Triple Exp. - Additive"] = [HWES3_ADD.sse] + [HWES3_ADD.aic] + [HWES3_ADD.bic]
results["Triple Exp. - Multiplicative"] = [HWES3_MUL.sse] + [HWES3_MUL.aic] + [HWES3_MUL.b]
# El Triple Exponencial Multiplicativo resalta por tener menor error SSE en comparación a
```

```
results
```

	Simple Exponencial	Double Exp. - Additive	Double Exp. - Multiplicative	Triple Exp. - Additive	Triple Exp. - Multiplicative
<b>SSE</b>	6.299774e+10	6.532978e+10	6.992348e+10	1.556908e+09	7.267883e+08
<b>AIC</b>	1.250321e+03	1.256502e+03	1.260579e+03	1.056297e+03	1.010588e+03
<b>BIC</b>	1.254510e+03	1.264880e+03	1.268957e+03	1.089807e+03	1.044098e+03

```
pd.concat([Sales_mes, HWES3_MUL_fitted], axis=1).plot(title='Original and fitted ts');
```



```
HWES3_MUL.summary()
```

ExponentialSmoothing Model Results

<b>Dep. Variable:</b>	<b>sales</b>	<b>No. Observations:</b>	60
<b>Model:</b>	ExponentialSmoothing	<b>SSE</b>	726788287.742
<b>Optimized:</b>	True	<b>AIC</b>	1010.588
<b>Trend:</b>	Multiplicative	<b>BIC</b>	1044.098
<b>Seasonal:</b>	Multiplicative	<b>AICC</b>	1027.271
<b>Seasonal Periods:</b>	12	<b>Date:</b>	Sun, 16 Apr 2023
<b>Box-Cox:</b>	False	<b>Time:</b>	04:06:25
<b>Box-Cox Coeff.:</b>	None		

	<b>coeff</b>	<b>code</b>	<b>optimized</b>
<b>smoothing_level</b>	0.5707143	alpha	True
<b>smoothing_trend</b>	0.0815306	beta	True
<b>smoothing_seasonal</b>	0.0001	gamma	True
<b>initial_level</b>	1.9917e+05	l.0	True
<b>initial_trend</b>	1.0137598	b.0	True
<b>initial_seasons.0</b>	0.7183591	s.0	True
<b>initial_seasons.1</b>	0.7190422	s.1	True
<b>initial_seasons.2</b>	0.9394036	s.2	True
<b>initial_seasons.3</b>	1.0577767	s.3	True
<b>initial_seasons.4</b>	1.1636667	s.4	True
<b>initial_seasons.5</b>	1.1891510	s.5	True
<b>initial_seasons.6</b>	1.3025210	s.6	True
<b>initial_seasons.7</b>	1.1481031	s.7	True
<b>initial_seasons.8</b>	1.0252301	s.8	True
<b>initial_seasons.9</b>	0.9804144	s.9	True
<b>initial_seasons.10</b>	1.0165740	s.10	True
<b>initial_seasons.11</b>	0.7397581	s.11	True

## ▼ Holt-Winters

```
fit1 = ExponentialSmoothing(
    Sales_mes,
    seasonal_periods=12,
    trend="add",
    seasonal="add",
    initialization_method="estimated",
).fit()
```

```
fit2 = ExponentialSmoothing(
```

```
Sales_mes,
seasonal_periods=12,
trend="add",
seasonal="mul",
initialization_method="estimated",
).fit()
```

```
fit3 = ExponentialSmoothing(
Sales_mes,
seasonal_periods=12,
trend="mul",
seasonal="add",
initialization_method="estimated",
).fit()
```

```
fit4 = ExponentialSmoothing(
Sales_mes,
seasonal_periods=12,
trend="mul",
seasonal="mul",
initialization_method="estimated",
).fit()
```

```
results2 = pd.DataFrame(
index=["SSE", "AIC", "BIC"])
```

```
results2["Holt Trend-Add Seas-Add"] = [fit1.sse] + [fit1.aic] + [fit1.bic]
results2["Holt Trend-Add Seas-Mul"] = [fit2.sse] + [fit2.aic] + [fit2.bic]
results2["Holt Trend-Mul Seas-Add"] = [fit3.sse] + [fit3.aic] + [fit3.bic]
results2["Holt Trend-Mul Seas-Mul"] = [fit4.sse] + [fit4.aic] + [fit4.bic]
```

```
results2
```

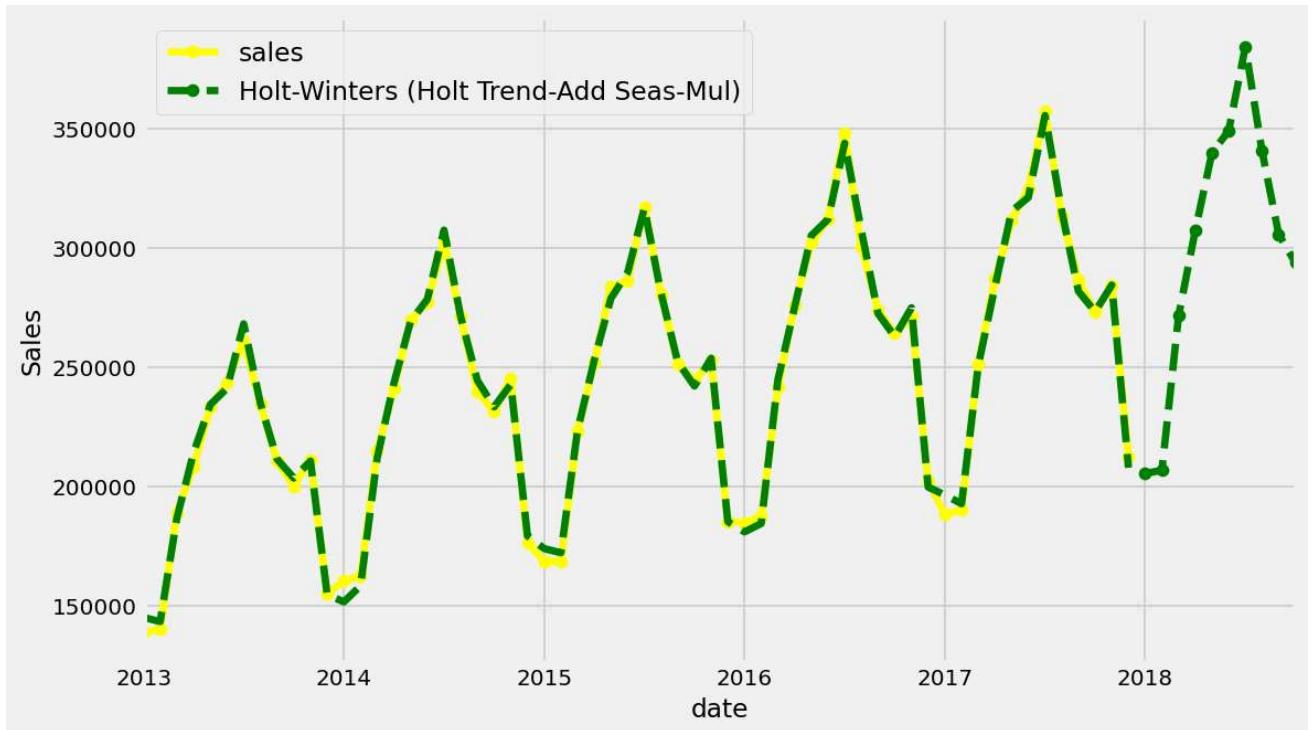
	Holt Trend-Add Seas-Add	Holt Trend-Add Seas-Mul	Holt Trend-Mul Seas-Add	Holt Trend-Mul Seas-Mul
<b>SSE</b>	1.556908e+09	7.054268e+08	1.600289e+09	7.267883e+08
<b>AIC</b>	1.056297e+03	1.008798e+03	1.057946e+03	1.010588e+03
<b>BIC</b>	1.089807e+03	1.042308e+03	1.091456e+03	1.044098e+03

Observamos que el mejor modelo lo otorga el '**Holt Trend-Add Seas-Mul**', dado que tiene un menor SSE que el Mul-Mul

## Predicción

```
ax = Sales_mes.plot(
figsize=(10, 6),
marker="o",
color="yellow",
```

```
)
ax.set_ylabel("Sales")
ax.set_xlabel("Year")
fit2.fittedvalues.plot(ax=ax, style="--", color="green")
fit2.forecast(10).rename("Holt-Winters (Holt Trend-Add Seas-Mul)").plot(
    ax=ax, style="--", marker="o", color="green", legend=True
)
plt.show()
# aqui obtamos solo por proyectar 10 periodos mensuales, dado que la base real tiene 60 pe
```



```
import numpy as np
from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error, median_abso

def evaluate_forecast(y,pred):
    results = pd.DataFrame({'r2_score':r2_score(y, pred),
```

```

        }, index=[0])
results['mean_absolute_error'] = mean_absolute_error(y, pred)
results['median_absolute_error'] = median_absolute_error(y, pred)
results['mse'] = mean_squared_error(y, pred)
results['msle'] = mean_squared_log_error(y, pred)
results['rmse'] = np.sqrt(results['mse'])
return results

```

```
evaluate_forecast(Sales_mes.sales, fit2.fittedvalues)
```

r2_score	mean_absolute_error	median_absolute_error	mse	msle
0.995726	0.15255	0.15255	0.00028	0.15255

## ▼ Auto - SARIMA

```

import warnings
warnings.filterwarnings("ignore")
def fxn():
    warnings.warn("deprecated", DeprecationWarning)
with warnings.catch_warnings():
    warnings.simplefilter("ignore")
    fxn()

import itertools
import numpy as np
import matplotlib.pyplot as plt
plt.style.use('fivethirtyeight')

import pandas as pd
pd.set_option('display.expand_frame_repr', False)
pd.set_option('display.max_columns', 500)
pd.set_option('display.width', 1000)

import statsmodels as sm
from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.graphics.tsaplots import plot_pacf
from statsmodels.tsa.stattools import adfuller
from statsmodels.tsa.seasonal import seasonal_decompose

from math import sqrt

import matplotlib
import matplotlib.pyplot as plt
matplotlib.rcParams['axes.labelsize'] = 14
matplotlib.rcParams['xtick.labelsize'] = 12
matplotlib.rcParams['ytick.labelsize'] = 12
matplotlib.rcParams['text.color'] = 'k'
import seaborn as sns

```

```
from random import random

from sklearn.metrics import mean_squared_error, r2_score, mean_absolute_error, median_abso

Sales_m = Sales_mes.rename(columns={'sales': 'y'})
Sales_m.head()
```

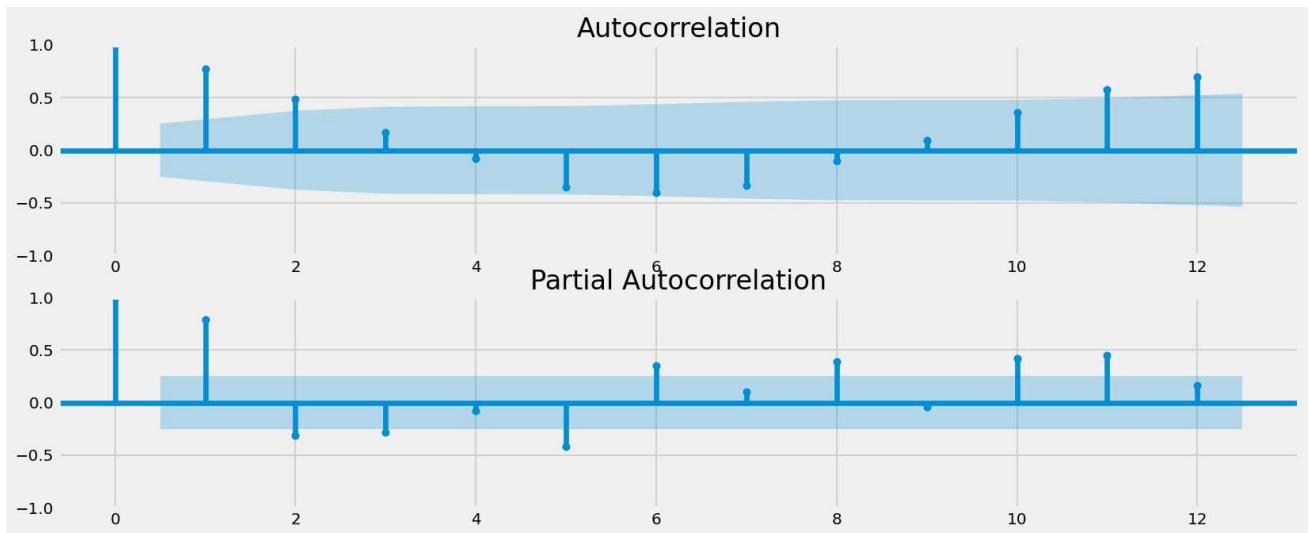
y 

	date
2013-01-31	138832.0
2013-02-28	140085.0
2013-03-31	189060.0
2013-04-30	208446.0
2013-05-31	233266.0

```
plt.figure(figsize=(15,6))
plt.plot(Sales_m)
plt.show()
```

```
from statsmodels.graphics.tsaplots import plot_acf
from statsmodels.graphics.tsaplots import plot_pacf

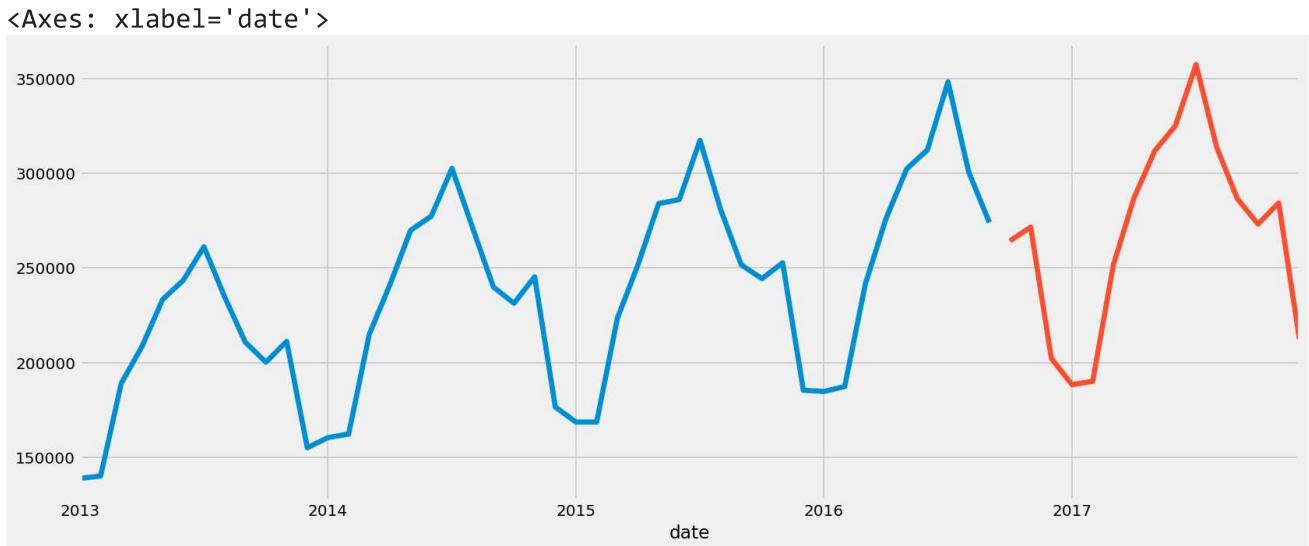
plt.figure(figsize=(15,6))
plt.subplot(211)
plot_acf(Sales_m, ax=plt.gca(), lags = 12)
plt.subplot(212)
plot_pacf(Sales_m, ax=plt.gca(), lags = 12)
plt.show()
```



```
#divide into train and validation set
train_m = Sales_m[:int(0.75*(len(Sales_m)))]
valid_m = Sales_m[int(0.75*(len(Sales_m))):]

#plotting the data
plt.figure(figsize=(15,6))
```

```
train_m['y'].plot()  
valid_m['y'].plot()
```



## Aplicando AUTO SARIMA para encontrar los mejores parámetros

```
pip install pmdarima
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/  
Requirement already satisfied: pmdarima in /usr/local/lib/python3.9/dist-packages (2  
Requirement already satisfied: scikit-learn>=0.22 in /usr/local/lib/python3.9/dist-pa  
Requirement already satisfied: scipy>=1.3.2 in /usr/local/lib/python3.9/dist-packages  
Requirement already satisfied: setuptools!=50.0.0,>=38.6.0 in /usr/local/lib/python3  
Requirement already satisfied: numpy>=1.21.2 in /usr/local/lib/python3.9/dist-package  
Requirement already satisfied: Cython!=0.29.18,!>=0.29.31,>=0.29 in /usr/local/lib/pyt  
Requirement already satisfied: pandas>=0.19 in /usr/local/lib/python3.9/dist-packages  
Requirement already satisfied: statsmodels>=0.13.2 in /usr/local/lib/python3.9/dist-p  
Requirement already satisfied: urllib3 in /usr/local/lib/python3.9/dist-packages (fr  
Requirement already satisfied: joblib>=0.11 in /usr/local/lib/python3.9/dist-packages  
Requirement already satisfied: python-dateutil>=2.8.1 in /usr/local/lib/python3.9/dis
```

```
Requirement already satisfied: pytz>=2020.1 in /usr/local/lib/python3.9/dist-packages
Requirement already satisfied: threadpoolctl>=2.0.0 in /usr/local/lib/python3.9/dist-
Requirement already satisfied: patsy>=0.5.2 in /usr/local/lib/python3.9/dist-packages
Requirement already satisfied: packaging>=21.3 in /usr/local/lib/python3.9/dist-pac
Requirement already satisfied: six in /usr/local/lib/python3.9/dist-packages (from pa
```

```
#building the model
from pmдарима import auto_arima
model = auto_arima(train_m, trace=True, error_action='ignore', suppress_warnings=True, sea
model.fit(train_m)
```

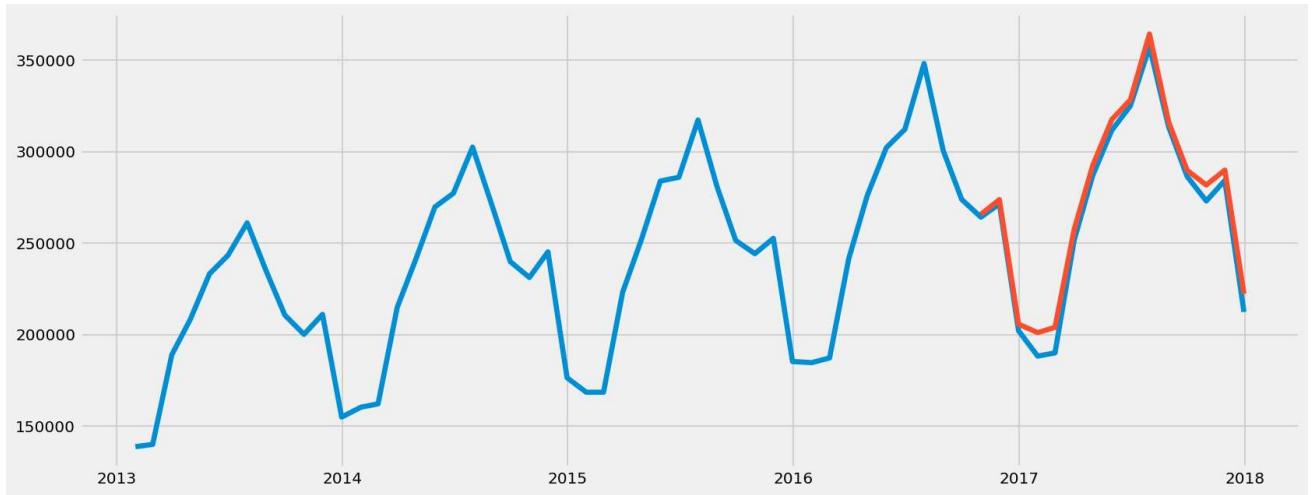
```
Performing stepwise search to minimize bic
start_index = valid_m.index.min()
end_index = valid_m.index.max()

#Predictions
pred = model.predict()

pred = model.predict(n_periods=len(valid))
pred = pd.DataFrame(pred,index = valid.index,columns=[ 'Prediction'])

forecast = model.predict(n_periods=len(valid))
forecast = pd.DataFrame(forecast,index = valid.index,columns=[ 'Prediction'])

plt.figure(figsize=(15,6))
#plot the predictions for validation set
plt.plot(Sales_m.y, label='Train')
#plt.plot(valid, label='Valid')
plt.plot(forecast, label='Prediction')
plt.show()
```

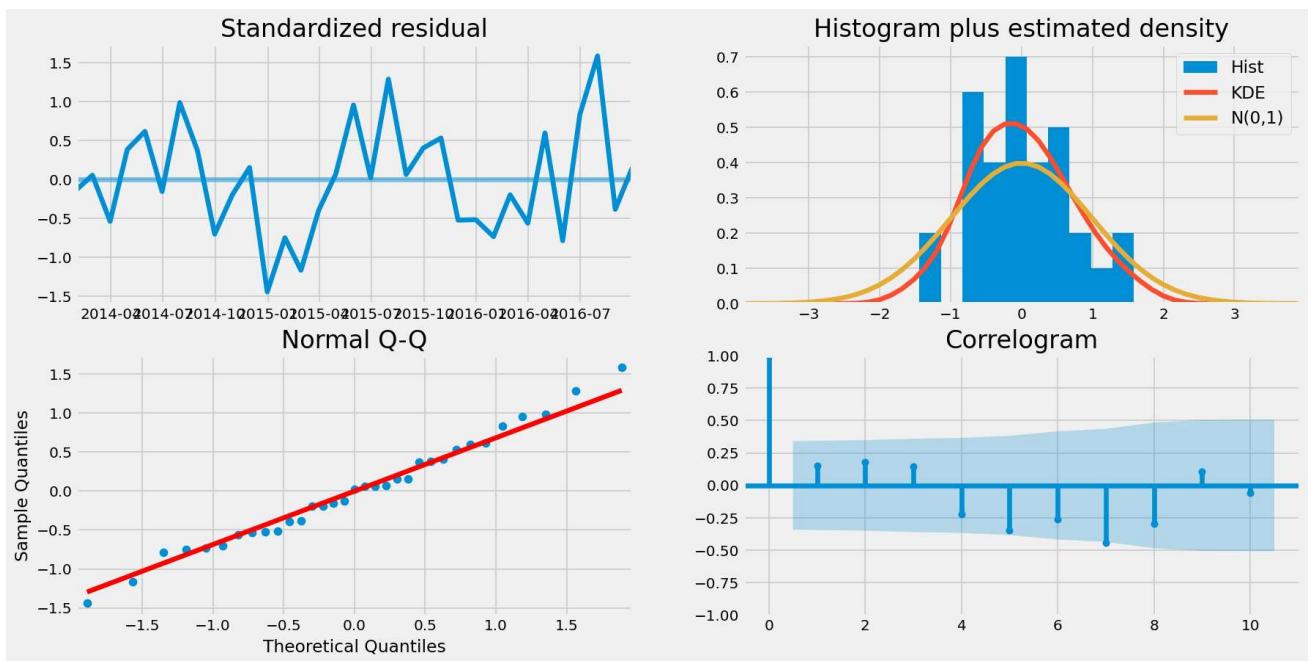


```
evaluate_forecast(Sales_m[start_index:end_index], forecast)
```

r2_score	mean_absolute_error	median_absolute_error	mse	msle
----------	---------------------	-----------------------	-----	------

0.070014	6132.773159	5682.701337	5.0862826407	0.00101	7131.8
----------	-------------	-------------	--------------	---------	--------

```
model.plot_diagnostics(figsize=(16, 8))
plt.show()
```



Diagnóstico del modelo:

\*Nuestra principal preocupación es garantizar que los residuos de nuestro modelo no estén correlacionados y normalmente se distribuyan con media cero.

\*Si el modelo estacional ARIMA no satisface estas propiedades, es una buena indicación de que puede mejorarse aún más.

El diagnóstico del modelo sugiere que el modelo residual se distribuye normalmente en función de lo siguiente:

\*En la gráfica superior derecha, la línea roja de KDE sigue de cerca con la línea N (0,1). Donde, N (0,1) es la notación estándar para una distribución normal con media 0 y desviación estándar de 1. Esta es una buena indicación de que los residuos se distribuyen normalmente.

\*La gráfica qq en la parte inferior izquierda muestra que la distribución ordenada de los residuos (puntos azules) sigue la tendencia lineal de las muestras tomadas de una distribución normal estándar. Nuevamente, esta es una fuerte indicación de que los residuos se distribuyen normalmente.

\*Los residuos a lo largo del tiempo (gráfico superior izquierdo) no muestran ninguna estacionalidad obvia y parecen ser ruido blanco.

\*Esto se confirma mediante el gráfico de autocorrelación (es decir, correlograma) en la parte inferior derecha, que muestra que los residuos de series temporales tienen una baja correlación con versiones rezagadas de sí mismo. Quitando el 0 que naturalmente es el mismo dato, y un dato que se escapa apenas del rango que debe ser por la variación que no es constante.

## ▼ Comentarios finales

El mejor modelo obtenido en el análisis es el Holt-Winters (Holt Trend-Add Seas-Mul) con un R2 de 99.6% y un RMSE de 3'428.9 (casi la mitad de lo obtenido en el Auto Sarima)

- A continuación los 10 períodos resultado del mejor modelo obtenido y su gráfico.

```
fit2.forecast(10)
```

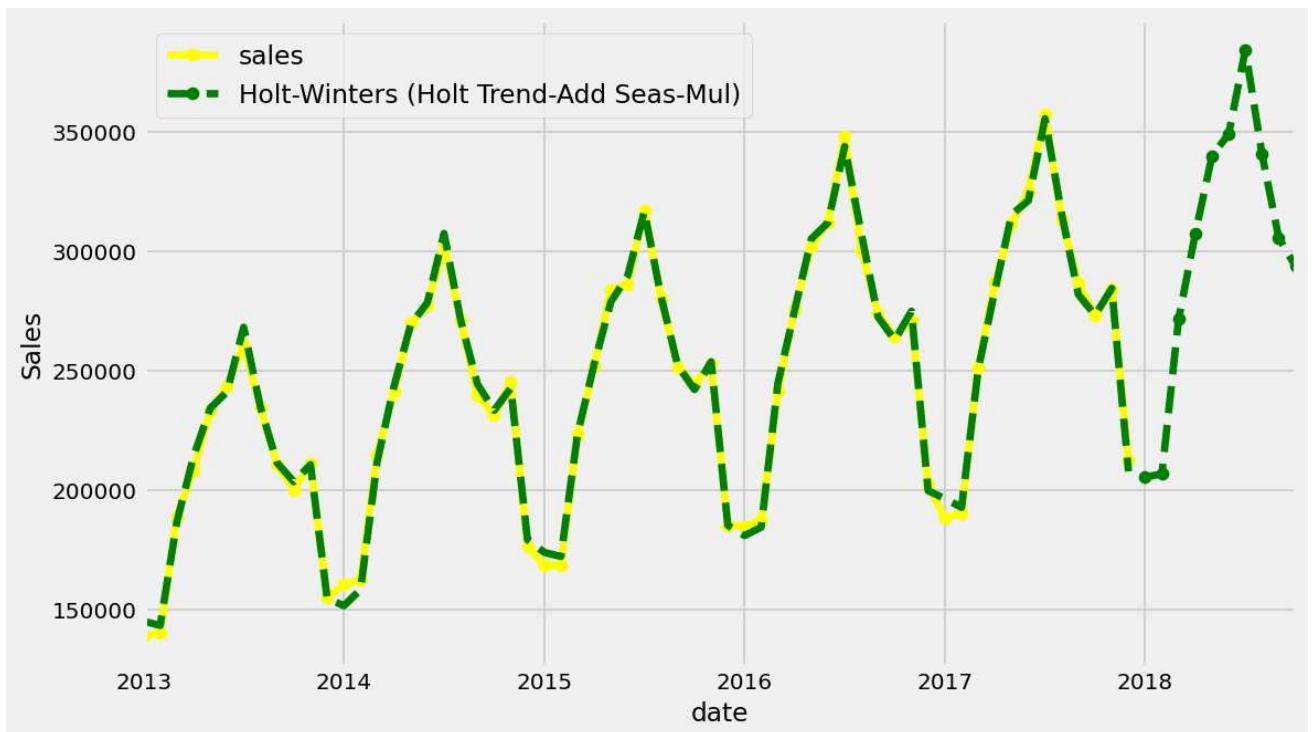
2018-01-31	205481.374898
2018-02-28	206768.957393
2018-03-31	271563.688341
2018-04-30	307388.001497
2018-05-31	339927.085304
2018-06-30	349177.485214
2018-07-31	384443.701129
2018-08-31	340609.768508
2018-09-30	305714.398009

```
2018-10-31      293839.151553
Freq: M, dtype: float64
```

```
ax = Sales_mes.plot(
    figsize=(10, 6),
    marker="o",
    color="yellow",
)

ax.set_ylabel("Sales")
ax.set_xlabel("Year")
fit2.fittedvalues.plot(ax=ax, style="--", color="green")
fit2.forecast(10).rename("Holt-Winters (Holt Trend-Add Seas-Mul)").plot(
    ax=ax, style="--", marker="o", color="green", legend=True
)

plt.show()
```





✓ 0 s completado a las 23:08

